

# Advanced Rust Programming

## Expert-Level Techniques and Internals

Rust Learning Series - Advanced Track

December 14, 2025



# Prerequisites

This book assumes mastery of all intermediate Rust concepts. You should have extensive practical experience with:

- **All Intermediate Topics:** Advanced traits, complex lifetimes, and macro programming
- **Unsafe Rust and FFI:** Comfortable writing unsafe code and interfacing with C
- **Advanced Async Programming:** Understanding futures, pinning, and async runtimes
- **Performance Optimization:** Experience with profiling and optimization techniques
- **Lock-Free Concurrency:** Understanding atomic operations and memory orderings

This book covers expert-level topics including compiler internals, advanced procedural macros, custom allocators, embedded systems programming, WebAssembly, async runtime implementation, advanced lock-free algorithms, language tooling development, performance profiling, and contributing to Rust itself.



# Contents

<b>Prerequisites</b>	<b>3</b>
<b>1 Compiler Internals and MIR</b>	<b>7</b>
1.1 The Rust Compilation Pipeline . . . . .	7
1.2 Working with MIR . . . . .	7
1.2.1 MIR Structure . . . . .	8
1.3 Compiler Plugins and Custom Lints . . . . .	8
<b>2 Advanced Procedural Macros</b>	<b>9</b>
2.1 Complex Derive Macros . . . . .	9
2.2 Parsing Custom Syntax . . . . .	9
2.3 Hygiene and Span Management . . . . .	10
<b>3 Custom Allocators</b>	<b>11</b>
3.1 The GlobalAlloc Trait . . . . .	11
3.2 Arena Allocators . . . . .	11
3.3 Per-Collection Allocators . . . . .	12
<b>4 No-Std and Embedded Systems</b>	<b>13</b>
4.1 No-Std Fundamentals . . . . .	13
4.2 Embedded HAL . . . . .	13
4.3 Volatile Memory Access . . . . .	14
4.4 Interrupt Handling . . . . .	14
<b>5 WebAssembly Deep Dive</b>	<b>15</b>
5.1 Wasm Bindgen . . . . .	15
5.2 JavaScript Interoperability . . . . .	15
5.3 Optimizing WebAssembly Size . . . . .	16
<b>6 Async Runtime Internals</b>	<b>17</b>
6.1 Building a Simple Executor . . . . .	17
6.2 Waker Implementation . . . . .	17
6.3 Reactor Pattern . . . . .	18
<b>7 Advanced Lock-Free Algorithms</b>	<b>19</b>
7.1 The ABA Problem . . . . .	19
7.2 Tagged Pointers . . . . .	19
7.3 Epoch-Based Reclamation . . . . .	20

<b>8 Building Language Tooling</b>	<b>21</b>
8.1 Using rust-analyzer APIs . . . . .	21
8.2 Custom Cargo Subcommands . . . . .	21
8.3 LSP Server Implementation . . . . .	22
<b>9 Performance Profiling and Optimization</b>	<b>23</b>
9.1 Benchmarking with Criterion . . . . .	23
9.2 CPU Profiling . . . . .	23
9.2.1 Profile-Guided Optimization . . . . .	24
9.3 Memory Profiling . . . . .	24
<b>10 Contributing to Rust Itself</b>	<b>25</b>
10.1 Rust Compiler Development . . . . .	25
10.2 The RFC Process . . . . .	25
10.3 Writing Compiler Tests . . . . .	26
<b>11 Summary and Mastery Path</b>	<b>27</b>
11.1 Key Takeaways . . . . .	27
11.2 Mastery Projects . . . . .	27
11.3 Advanced Resources . . . . .	28
11.3.1 Books and Documentation . . . . .	28
11.3.2 Research and Papers . . . . .	28
11.3.3 Communities . . . . .	28
11.4 The Complete Journey . . . . .	28
11.5 Conclusion . . . . .	28

# Chapter 1

## Compiler Internals and MIR

Understanding Rust's compilation pipeline provides insights into how the language works and enables advanced compiler development.

### 1.1 The Rust Compilation Pipeline

The Rust compiler (rustc) transforms source code into machine code through several intermediate representations:

1. **Lexing and Parsing:** Source code is tokenized and parsed into an Abstract Syntax Tree (AST)
2. **Macro Expansion:** Macros are expanded to generate more AST nodes
3. **HIR (High-level IR):** AST is lowered to a more compiler-friendly representation
4. **Type Checking and Inference:** Type information is inferred and checked
5. **MIR (Mid-level IR):** HIR is lowered to MIR for borrow checking and optimization
6. **Borrow Checking:** MIR is analyzed to enforce borrowing rules
7. **Optimization:** Various optimizations are applied to MIR
8. **LLVM IR:** MIR is translated to LLVM's intermediate representation
9. **Machine Code:** LLVM generates architecture-specific machine code

### 1.2 Working with MIR

Mid-level Intermediate Representation (MIR) is crucial for borrow checking and optimization. You can inspect MIR for your functions:

```
1 // View MIR for a function
2 #[rustc_dump_mir(before = "all", after = "all")]
3 fn example(x: i32) -> i32 {
4     x * 2 + 1
5 }
```

### 1.2.1 MIR Structure

MIR represents functions as control flow graphs with basic blocks:

```

1 // Simplified MIR representation
2 fn example(_1: i32) -> i32 {
3     let mut _0: i32;
4     let mut _2: i32;
5     let mut _3: i32;
6
7     bb0: {
8         _2 = _1;
9         _3 = const 2_i32;
10        _2 = Mul(move _2, move _3);
11        _3 = const 1_i32;
12        _0 = Add(move _2, move _3);
13        return;
14    }
15 }
```

MIR is used for:

- **Borrow Checking:** Ensuring ownership and borrowing rules
- **Optimization:** Performing transformations that preserve semantics
- **Const Evaluation:** Computing compile-time constants

## 1.3 Compiler Plugins and Custom Lints

You can extend rustc with custom lints:

```

1 #![feature(rustc_private)]
2 extern crate rustc_lint;
3 extern crate rustc_middle;
4
5 use rustc_lint::{LateContext, LateLintPass, LintContext};
6
7 declare_lint! {
8     pub CUSTOM_LINT,
9     Warn,
10    "description of custom lint"
11 }
12
13 impl<'tcx> LateLintPass<'tcx> for CustomLint {
14     fn check_expr(&mut self, cx: &LateContext<'tcx>, expr: &Expr) {
15         // Custom lint logic
16         // Analyze expressions and emit warnings
17     }
18 }
```

Custom lints enable enforcement of project-specific coding standards and detection of anti-patterns.

# Chapter 2

## Advanced Procedural Macros

Procedural macros enable powerful compile-time metaprogramming. Advanced techniques allow you to create complex code generation.

### 2.1 Complex Derive Macros

Derive macros automatically implement traits. Building sophisticated derive macros requires parsing attributes and generating complex code:

```
1 #[proc_macro_derive(Builder, attributes(builder))]
2 pub fn derive_builder(input: TokenStream) -> TokenStream {
3     let input = parse_macro_input!(input as DeriveInput);
4
5     let name = &input.ident;
6     let builder_name = format_ident!("{}Builder", name);
7
8     let fields = match &input.data {
9         Data::Struct(data) => &data.fields,
10        _ => panic!("Builder only works on structs"),
11    };
12
13     // Generate builder methods for each field
14     let methods = fields.iter().map(|f| {
15         let field_name = &f.ident;
16         let field_type = &f.ty;
17         quote! {
18             pub fn #field_name(mut self, value: #field_type) -> Self {
19                 self.#field_name = Some(value);
20                 self
21             }
22         }
23     });
24
25     // Generate the builder struct and implementation
26     // ... more complex logic
27 }
```

### 2.2 Parsing Custom Syntax

The `syn` crate enables parsing arbitrary syntax:

```
1 use syn::parse::Parse, Token;
2
3 struct MyMacroInput {
4     name: Ident,
```

```

5     _arrow: Token! [=>] ,
6     value: Expr ,
7 }
8
9 impl Parse for MyMacroInput {
10    fn parse(input: ParseStream) -> Result<Self> {
11        Ok(MyMacroInput {
12            name: input.parse()?,
13            _arrow: input.parse()?,
14            value: input.parse()?,
15        })
16    }
17 }
18
19 #[proc_macro]
20 pub fn my_macro(input: TokenStream) -> TokenStream {
21     let parsed = parse_macro_input!(input as MyMacroInput);
22     // Use parsed.name and parsed.value to generate code
23 }
```

## 2.3 Hygiene and Span Management

Macro hygiene prevents identifier conflicts. Span management enables precise error reporting:

```

1 use proc_macro::Span;
2 use quote::quote_spanned;
3
4 fn generate_with_span(span: Span) -> TokenStream {
5     quote_spanned! {span=>
6         compile_error!("Error at original location");
7     }.into()
8 }
9
10 // Preserve spans for better error messages
11 let tokens = quote_spanned! {field.span()=>
12     self.#field_name = value;
13 };
```

Macro-generated identifiers automatically avoid conflicts with user code, but you can control scoping when necessary.

# Chapter 3

## Custom Allocators

Custom allocators provide control over memory management for specialized use cases.

### 3.1 The GlobalAlloc Trait

Implement custom allocators using `GlobalAlloc`:

```
1 use std::alloc::{GlobalAlloc, Layout, System};
2
3 struct MyAllocator;
4
5 unsafe impl GlobalAlloc for MyAllocator {
6     unsafe fn alloc(&self, layout: Layout) -> *mut u8 {
7         println!("Allocating {} bytes", layout.size());
8         System.alloc(layout)
9     }
10
11    unsafe fn dealloc(&self, ptr: *mut u8, layout: Layout) {
12        println!("Deallocating {} bytes", layout.size());
13        System.dealloc(ptr, layout)
14    }
15
16 #[global_allocator]
17 static ALLOCATOR: MyAllocator = MyAllocator;
```

### 3.2 Arena Allocators

Arena allocators provide fast bulk allocation and deallocation:

```
1 struct Arena {
2     buf: Vec<u8>,
3     offset: usize,
4 }
5
6 impl Arena {
7     fn alloc<T>(&mut self, value: T) -> &mut T {
8         let layout = Layout::new::<T>();
9         let offset = self.offset;
10        self.offset += layout.size();
11
12        unsafe {
13            let ptr = self.buf.as_mut_ptr().add(offset) as *mut T;
14            ptr.write(value);
15            &mut *ptr
16        }
17    }
18 }
```

```

16     }
17 }
18
19 fn reset(&mut self) {
20     self.offset = 0;
21     // No individual drops - bulk deallocation
22 }
23 }
```

Arena allocators excel when you allocate many objects with similar lifetimes and deallocate them all at once.

### 3.3 Per-Collection Allocators

The `allocator_api` feature enables per-collection custom allocators:

```

1 #[feature(allocator_api)]
2
3 use std::alloc::Allocator;
4
5 struct BumpAllocator { /* ... */ }
6
7 unsafe impl Allocator for BumpAllocator {
8     fn allocate(&self, layout: Layout)
9         -> Result<NonNull<[u8]>, AllocError>
10    {
11        // Bump allocation logic
12    }
13
14     unsafe fn deallocate(&self, ptr: NonNull<u8>, layout: Layout) {
15         // No-op for bump allocator
16     }
17 }
18
19 // Use with collections
20 let vec: Vec<i32, BumpAllocator> = Vec::new_in(bump_allocator);
```

# Chapter 4

## No-Std and Embedded Systems

Rust's no-std mode enables bare-metal programming for embedded systems and operating system kernels.

### 4.1 No-Std Fundamentals

In no-std environments, you don't have access to the standard library:

```
1 #![no_std]
2 #![no_main]
3
4 use core::panic::PanicInfo;
5
6 #[panic_handler]
7 fn panic(_info: &PanicInfo) -> ! {
8     loop {}
9 }
10
11 #[no_mangle]
12 pub extern "C" fn _start() -> ! {
13     // Entry point for bare metal
14     loop {}
15 }
16
17 // Use core instead of std
18 use core::ptr;
19 use core::mem;
```

No-std environments have no heap allocation, no operating system services, and minimal runtime support.

### 4.2 Embedded HAL

The Hardware Abstraction Layer enables portable embedded code:

```
1 use embedded_hal::digital::v2::OutputPin;
2
3 struct Led<P: OutputPin> {
4     pin: P,
5 }
6
7 impl<P: OutputPin> Led<P> {
8     fn on(&mut self) -> Result<(), P::Error> {
9         self.pin.set_high()
10    }
11 }
```

```

12     fn off(&mut self) -> Result<(), P::Error> {
13         self.pin.set_low()
14     }
15 }
16
17 // Works with any GPIO implementation
18 let mut led = Led { pin: gpio_pin };
19 led.on().unwrap();

```

### 4.3 Volatile Memory Access

Memory-mapped I/O requires volatile operations:

```

1 use core::ptr::{read_volatile, write_volatile};
2
3 // Memory-mapped IO
4 const GPIO_BASE: usize = 0x4000_0000;
5
6 #[repr(C)]
7 struct GpioRegisters {
8     data: u32,
9     direction: u32,
10    interrupt: u32,
11 }
12
13 fn set_gpio(bit: u8) {
14     unsafe {
15         let gpio = GPIO_BASE as *mut GpioRegisters;
16         let mut data = read_volatile(&(*gpio).data);
17         data |= 1 << bit;
18         write_volatile(&mut (*gpio).data, data);
19     }
20 }

```

Volatile operations prevent compiler optimizations that would be incorrect for hardware registers.

### 4.4 Interrupt Handling

Embedded systems use interrupts for asynchronous events:

```

1 use cortex_m_rt::interrupt;
2
3 #[interrupt]
4 fn TIM2() {
5     // Timer 2 interrupt handler
6     static mut COUNT: u32 = 0;
7
8     unsafe {
9         *COUNT += 1;
10        // Clear interrupt flag
11        (*TIM2::ptr()).sr.modify(|_, w| w.uif().clear_bit());
12    }
13 }
14
15 // Configure interrupt
16 unsafe {
17     cortex_m::peripheral::NVIC::unmask(Interrupt::TIM2);
18 }

```

# Chapter 5

## WebAssembly Deep Dive

WebAssembly enables running Rust code in web browsers with near-native performance.

### 5.1 Wasm Bindgen

The `wasm-bindgen` tool generates JavaScript bindings:

```
1 use wasm_bindgen::prelude::*;

2 #[wasm_bindgen]
3 pub fn fibonacci(n: u32) -> u32 {
4     match n {
5         0 => 0,
6         1 => 1,
7         _ => fibonacci(n - 1) + fibonacci(n - 2),
8     }
9 }

10 #[wasm_bindgen]
11 extern "C" {
12     #[wasm_bindgen(js_namespace = console)]
13     fn log(s: &str);
14 }
15

16 #[wasm_bindgen]
17 pub fn greet(name: &str) {
18     log(&format!("Hello, {}!", name));
19 }
20

21 }
```

### 5.2 JavaScript Interoperability

The `web-sys` crate provides Web API bindings:

```
1 use wasm_bindgen::JsCast;
2 use web_sys::{Document, Element, HtmlElement};

3 #[wasm_bindgen(start)]
4 pub fn main() -> Result<(), JsValue> {
5     let window = web_sys::window().unwrap();
6     let document = window.document().unwrap();

7     let body = document.body().unwrap();
8     let div = document.create_element("div")?;
9     div.set_inner_html("Hello from Rust!");

10    body.append_child(div).unwrap();
11 }

12 }
```

```
13     body.append_child(&div)?;  
14     Ok(())  
15 }
```

### 5.3 Optimizing WebAssembly Size

Size optimization is crucial for web delivery:

```
1 # Cargo.toml  
2 [profile.release]  
3 opt-level = "z"          # Optimize for size  
4 lto = true                # Link-time optimization  
5 codegen-units = 1         # Better optimization  
6 panic = "abort"           # Smaller binary  
7 strip = true              # Strip symbols
```

Additional techniques:

- Avoid formatting macros in release builds
- Use `wee_alloc` for smaller allocator
- Tree-shake with `wasm-gc`
- Post-process with `wasm-opt`

# Chapter 6

## Async Runtime Internals

Building custom async runtimes provides deep understanding of how async works in Rust.

### 6.1 Building a Simple Executor

An executor polls futures to completion:

```
1 use std::future::Future;
2 use std::task::{Context, Poll, RawWaker, RawWakerVTable, Waker};
3
4 struct SimpleExecutor {
5     tasks: Vec<Pin<Box<dyn Future<Output = ()>>>,
6 }
7
8 impl SimpleExecutor {
9     fn run(&mut self) {
10         while !self.tasks.is_empty() {
11             self.tasks.retain_mut(|task| {
12                 let waker = create_waker();
13                 let mut cx = Context::from_waker(&waker);
14
15                 match task.as_mut().poll(&mut cx) {
16                     Poll::Ready(_) => false, // Remove
17                     Poll::Pending => true, // Keep
18                 }
19             });
20         }
21     }
22 }
```

### 6.2 Waker Implementation

Wakers notify the executor when futures are ready:

```
1 fn create_waker() -> Waker {
2     unsafe fn clone(ptr: *const ()) -> RawWaker {
3         RawWaker::new(ptr, &VTABLE)
4     }
5
6     unsafe fn wake(_: *const ()) {
7         // Wake the task
8     }
9
10    unsafe fn wake_by_ref(_: *const ()) {
11        // Wake without consuming
```

```

12 }
13
14 unsafe fn drop(_: *const ()) {}
15
16 static VTABLE: RawWakerVTable = RawWakerVTable::new(
17     clone, wake, wake_by_ref, drop
18 );
19
20 let raw = RawWaker::new(std::ptr::null(), &VTABLE);
21 unsafe { Waker::from_raw(raw) }
22 }
```

### 6.3 Reactor Pattern

Reactors handle I/O events:

```

1 use mio::{Events, Interest, Poll, Token};
2
3 struct Reactor {
4     poll: Poll,
5     events: Events,
6     handlers: HashMap<Token, Box<dyn FnMut()>>,
7 }
8
9 impl Reactor {
10     fn register<S: Source>(&mut self, source: &mut S,
11                             handler: impl FnMut() + 'static)
12     {
13         let token = Token(self.handlers.len());
14         self.poll.registry()
15             .register(source, token, Interest::READABLE)
16             .unwrap();
17         self.handlers.insert(token, Box::new(handler));
18     }
19
20     fn run(&mut self) {
21         loop {
22             self.poll.poll(&mut self.events, None).unwrap();
23             for event in &self.events {
24                 if let Some(handler) = self.handlers.get_mut(&event.token()) {
25                     handler();
26                 }
27             }
28         }
29     }
30 }
```

# Chapter 7

## Advanced Lock-Free Algorithms

Lock-free data structures enable high-performance concurrent programming without locks.

### 7.1 The ABA Problem

Lock-free algorithms using compare-and-swap can suffer from the ABA problem:

```
1 use std::sync::atomic::{AtomicPtr, AtomicUsize, Ordering};
2
3 struct Node<T> {
4     data: T,
5     next: *mut Node<T>,
6 }
7
8 struct Stack<T> {
9     head: AtomicPtr<Node<T>>,
10 }
11
12 impl<T> Stack<T> {
13     fn push(&self, data: T) {
14         let node = Box::into_raw(Box::new(Node {
15             data,
16             next: std::ptr::null_mut(),
17         }));
18
19         loop {
20             let head = self.head.load(Ordering::Acquire);
21             unsafe { (*node).next = head; }
22
23             if self.head.compare_exchange(
24                 head, node,
25                 Ordering::Release, Ordering::Acquire
26             ).is_ok() {
27                 break;
28             }
29         }
30     }
31 }
```

### 7.2 Tagged Pointers

Tagged pointers solve the ABA problem by including a version counter:

```
1 // Pack counter with pointer
2 struct Tagged<T> {
```

```

3     ptr: usize, // Bottom bits: counter, top bits: pointer
4 }
5
6 impl<T> Tagged<T> {
7     fn new(ptr: *mut T, tag: usize) -> Self {
8         let addr = ptr as usize;
9         Tagged { ptr: addr | (tag & 0xFFFF) }
10    }
11
12    fn get_ptr(&self) -> *mut T {
13        (self.ptr & !0xFFFF) as *mut T
14    }
15
16    fn get_tag(&self) -> usize {
17        self.ptr & 0xFFFF
18    }
19}
20
21 // Use AtomicUsize for tagged pointer
22 struct AbaFreeStack<T> {
23     head: AtomicUsize,
24 }
```

### 7.3 Epoch-Based Reclamation

The `crossbeam-epoch` crate provides safe memory reclamation:

```

1 use crossbeam_epoch::{self as epoch, Atomic, Owned};
2
3 struct Node<T> {
4     data: T,
5     next: Atomic<Node<T>>,
6 }
7
8 struct Stack<T> {
9     head: Atomic<Node<T>>,
10 }
11
12 impl<T> Stack<T> {
13     fn push(&self, data: T) {
14         let node = Owned::new(Node {
15             data,
16             next: Atomic::null(),
17         });
18
19         let guard = epoch::pin();
20         loop {
21             let head = self.head.load(Ordering::Acquire, &guard);
22             node.next.store(head, Ordering::Relaxed);
23
24             if self.head.compare_exchange(
25                 head, node,
26                 Ordering::Release, Ordering::Acquire, &guard
27             ).is_ok() {
28                 break;
29             }
30         }
31     }
32 }
```

# Chapter 8

## Building Language Tooling

Creating tools for Rust development enhances the ecosystem and provides deep language understanding.

### 8.1 Using rust-analyzer APIs

The `rust-analyzer` APIs enable code analysis:

```
1 use ra_ap_syntax::{ast, AstNode, SyntaxKind};
2 use ra_ap_ide::{Analysis, AnalysisHost, FileId};
3
4 fn analyze_code(source: &str) -> Vec<String> {
5     let parse = ast::SourceFile::parse(source);
6     let root = parse.tree();
7
8     let mut functions = Vec::new();
9
10    for node in root.syntax().descendants() {
11        if let Some(func) = ast::Fn::cast(node) {
12            if let Some(name) = func.name() {
13                functions.push(name.to_string());
14            }
15        }
16    }
17
18    functions
19 }
```

### 8.2 Custom Cargo Subcommands

Cargo subcommands extend Cargo's functionality:

```
1 // cargo-mycmd/src/main.rs
2 use clap::Parser;
3
4 #[derive(Parser)]
5 #[command(name = "cargo")]
6 #[command(bin_name = "cargo")]
7 enum Cargo {
8     Mycmd(Args),
9 }
10
11 #[derive(Parser)]
12 struct Args {
13     #[arg(long)]
```

```

14     verbose: bool,
15 }
16
17 fn main() {
18     let Cargo::Mycmd(args) = Cargo::parse();
19
20     // Access cargo metadata
21     let metadata = cargo_metadata::MetadataCommand::new()
22         .exec()
23         .unwrap();
24
25     // Implement custom logic
26 }
```

Install as `cargo-mycmd` and invoke with `cargo mycmd`.

### 8.3 LSP Server Implementation

Language Server Protocol enables IDE integration:

```

1 use tower_lsp::{LspService, Server};
2 use tower_lsp::lsp_types::*;
3
4 struct Backend;
5
6 #[tower_lsp::async_trait]
7 impl LanguageServer for Backend {
8     async fn initialize(&self, _: InitializeParams)
9         -> Result<InitializeResult>
10    {
11        Ok(InitializeResult {
12            capabilities: ServerCapabilities {
13                text_document_sync: Some(
14                    TextDocumentSyncCapability::Kind(
15                        TextDocumentSyncKind::FULL
16                    )
17                ),
18                completion_provider: Some(CompletionOptions::default()),
19                ..Default::default()
20            },
21            ..Default::default()
22        })
23    }
24
25    async fn completion(&self, _: CompletionParams)
26        -> Result<Option<CompletionResponse>>
27    {
28        // Provide completions
29    }
30 }
```

# Chapter 9

# Performance Profiling and Optimization

Profiling identifies bottlenecks and guides optimization efforts.

## 9.1 Benchmarking with Criterion

Criterion provides statistical benchmarking:

```
1 use criterion::{black_box, criterion_group, criterion_main, Criterion};
2
3 fn fibonacci_benchmark(c: &mut Criterion) {
4     c.bench_function("fib 20", |b| {
5         b.iter(|| fibonacci(black_box(20)))
6     });
7
8     c.bench_function("fib_iterative 20", |b| {
9         b.iter(|| fib_iterative(black_box(20)))
10    });
11 }
12
13 criterion_group!(benches, fibonacci_benchmark);
14 criterion_main!(benches);
```

Criterion features:

- Statistical analysis
- HTML reports with plots
- Comparison across runs

## 9.2 CPU Profiling

Various tools provide CPU profiling:

```
# Using perf on Linux
$ cargo build --release
$ perf record --call-graph=dwarf ./target/release/myapp
$ perf report

# Using flamegraph
$ cargo install flamegraph
```

```
$ cargo flamegraph

# Using samply (modern alternative)
$ cargo install samply
$ samply record ./target/release/myapp
```

### 9.2.1 Profile-Guided Optimization

PGO optimizes based on actual runtime profiles:

```
# 1. Build with instrumentation
RUSTFLAGS="-Cprofile-generate=/tmp/pgo" cargo build --release

# 2. Run typical workloads
./target/release/myapp < typical_input.txt

# 3. Rebuild with profile data
RUSTFLAGS="-Cprofile-use=/tmp/pgo" cargo build --release
```

## 9.3 Memory Profiling

Memory profiling identifies allocation patterns:

```
1 // Using dhat for heap profiling
2 #[global_allocator]
3 static ALLOC: dhat::Alloc = dhat::Alloc;
4
5 fn main() {
6     let _profiler = dhat::Profiler::new_heap();
7
8     // Your code here
9     let v: Vec<u64> = (0..1_000_000).collect();
10
11    // Profiler drops, generating report
12 }
```

Valgrind provides another approach:

```
$ valgrind --tool=massif ./target/release/myapp
$ ms_print massif.out.12345
```

# Chapter 10

## Contributing to Rust Itself

Contributing to Rust deepens your understanding and benefits the entire ecosystem.

### 10.1 Rust Compiler Development

Building rustc from source:

```
# Clone and build rustc
$ git clone https://github.com/rust-lang/rust.git
$ cd rust
$ ./x.py build

# Run tests
$ ./x.py test

# Build specific component
$ ./x.py build library/std

# Build documentation
$ ./x.py doc
```

Key areas for contribution:

- **Compiler:** Type checking, borrow checking, MIR
- **Standard Library:** Core, alloc, std
- **Cargo:** Build system and package manager
- **Rustdoc:** Documentation generator
- **Clippy:** Linter

### 10.2 The RFC Process

Rust uses RFCs (Requests for Comments) for significant changes:

1. **Idea:** Discuss on [internals.rust-lang.org](https://internals.rust-lang.org)
2. **RFC:** Submit RFC with detailed design
3. **Discussion:** Community reviews and suggests changes

4. **FCP:** Final Comment Period (10 days)
5. **Merge:** RFC accepted, implementation begins
6. **Implementation:** Write code, tests, documentation
7. **Stabilization:** Feature gate removal after testing

### 10.3 Writing Compiler Tests

Rustc uses comprehensive test suites:

```
1 // tests/ui/my-feature.rs
2 fn main() {
3     let x: i32 = "hello"; //~ ERROR mismatched types
4 }
```

Running tests:

```
# Run specific test
$ ./x.py test tests/ui/my-feature.rs

# Update expected output
$ ./x.py test tests/ui --bless
```

Test types:

- **ui:** Compiler error/warning tests
- **compile-fail:** Tests that should fail compilation
- **run-pass:** Tests that should compile and run
- **incremental:** Incremental compilation tests

# Chapter 11

## Summary and Mastery Path

### 11.1 Key Takeaways

This book has covered expert-level Rust topics:

1. **Compiler Internals:** Understanding MIR and the compilation pipeline
2. **Procedural Macros:** Mastering complex code generation
3. **Custom Allocators:** Implementing specialized memory management
4. **No-Std Programming:** Building for embedded and bare metal
5. **WebAssembly:** Deploying Rust to the web
6. **Async Runtimes:** Building custom executors and reactors
7. **Lock-Free Algorithms:** Implementing advanced concurrent data structures
8. **Language Tooling:** Creating tools for Rust development
9. **Performance Profiling:** Optimizing for maximum performance
10. **Contributing to Rust:** Giving back to the ecosystem

### 11.2 Mastery Projects

Challenge yourself with expert-level projects:

- Build a custom async runtime
- Implement a garbage collector
- Create a programming language in Rust
- Write an operating system kernel
- Build a database engine from scratch
- Implement a JIT compiler
- Create embedded firmware for real hardware
- Contribute features to rustc or cargo

## 11.3 Advanced Resources

### 11.3.1 Books and Documentation

- **Rustc Dev Guide:** [rustc-dev-guide.rust-lang.org](https://rustc-dev-guide.rust-lang.org)
- **Embedded Rust Book:** [docs.rust-embedded.org](https://docs.rust-embedded.org)
- **Rust and WebAssembly Book:** [rustwasm.github.io](https://rustwasm.github.io)
- **Rust Forge:** [forge.rust-lang.org](https://forge.rust-lang.org)

### 11.3.2 Research and Papers

- Lock-free algorithms and data structures
- Memory models and concurrency
- Type system theory
- Compiler optimization techniques

### 11.3.3 Communities

- **Zulip:** [rust-lang.zulipchat.com](https://rust-lang.zulipchat.com) - Real-time chat
- **Internals Forum:** [internals.rust-lang.org](https://internals.rust-lang.org) - Design discussions
- **GitHub:** [github.com/rust-lang](https://github.com/rust-lang) - Source code
- **Working Groups:** Async, embedded, WebAssembly, compiler
- **This Week in Rust:** Weekly newsletter

## 11.4 The Complete Journey

You've completed the full Rust learning path from beginner to expert:

- **Beginner:** Fundamentals of ownership, types, and basic patterns
- **Intermediate:** Advanced traits, lifetimes, macros, and async
- **Advanced:** Compiler internals, systems programming, and ecosystem contribution

## 11.5 Conclusion

Congratulations on reaching expert level! You're now equipped to:

- Build production systems in Rust
- Contribute to the Rust compiler and ecosystem
- Tackle the most challenging programming problems
- Mentor others in their Rust journey

The Rust community is welcoming and collaborative. Whether you're optimizing performance-critical code, building reliable systems, exploring language design, or teaching others, your expertise is valuable.

**Keep exploring, keep building, and keep contributing!**